




# Northeastern University



## CSYE 7220 DevOps Lecture 4 Docker Networks & Docker Compose

*23 May 2024*





Part 1

# DOCKER NETWORKS



# Creating a network for containers to share

## ❑ Communication issues? Create a network!

- `docker network create --driver bridge sa_network`
- `docker network inspect sa_network`

## ❑ Create a network with a specific subnet and gateway:

- `docker network create --driver=bridge --  
subnet=192.168.2.0/24 --gateway=192.168.2.10  
sa_network`

## ❑ Connect a container to a network

- `docker run --network= sa_network -itd --name=sa-fe-container sa-fe-image`
- Any other container you create on this network would be able to automatically connect to one another

## ❑ To remove a user-defined bridge network

- If containers are connected to the network, disconnect them first:  
`docker network disconnect network sa-fe-container`
- `docker network rm sa-network`

## ❑ <https://docs.docker.com/network/bridge/>

# A note about networks

- A network is *not* an IP!
- A network is like a city with many houses, *where each house has an IP*
- We create a network with docker network so that all the containers living in the network can talk to each other





# Containers communicating *through the host*

For Docker Desktop, replace <http://192.168.99.100> with <http://localhost>

## □ sa-logic:

- `docker build -t nsalo .`
- `docker run -p 5050:5000 nsalo`
- <http://192.168.99.100:5050/testHealth> works?
- <http://192.168.99.100:8080/testComms> works?



## □ sa-webapp:

- `mvn install`
- `docker build -t nsawa .`
- `docker run -p 8080:8080 -e "SA_LOGIC_API_URL=http://192.168.99.100:5050" nsawa`
- <http://192.168.99.100:8080/testHealth> works?
- <http://192.168.99.100:8080/testComms> works?



## □ sa-frontend:

- `npm install`
- `npm run build`
- `docker build -t nsafe .`
- `docker run -p 8085:80 nsafe`



- Does front end Send button work?: <http://192.168.99.100:8085>





## □ sa-logic: Containers communicating through a bridge network

- docker network create --driver bridge sanet
- docker build -t nsalo .
- docker run -p 5050:5000 --network=sanet nsalo
- docker network inspect sanet
  - Get the IP of the container: **a.b.c.d**
- <http://a.b.c.d:5000/testHealth> works?
- How about <http://192.168.99.100:5050/testComms> ?



## □ sa-webapp:

- mvn install
- docker build -t nsawa .
- docker run -p 8080:8080 --network=sanet -e "SA\_LOGIC\_API\_URL=http://**a.b.c.d**:5000" nsawa
- <http://p.q.r.s:8080/testHealth> works?
- How about <http://192.168.99.100:8080/testComms> ?



## □ sa-frontend:

- npm install && npm run build
- docker build -t nsafe .
- docker run -p 8085:80 --network=sanet nsafe
- docker network inspect



- Does the front end Send button work?: <http://192.168.99.100:8085>



# Results

- Don't forget to replace <http://192.168.99.100> with <http://localhost> if you're running on Docker Desktop!
- You should find that you cannot access any service endpoints at <http://a.b.c.d:5000/> or <http://p.q.r.s:8080/>, but your Send button at <http://192.168.99.100:8085> should still work (without having to redirect through your host)!
  - But keep accessing your service endpoint UI like [/testHealth](#) and [/testComms](#) through your host and port redirection instead!




# About networks

- In our native (non containerized) runs, we were using the *host network* as the network for all 3 apps to connect on
  - That's your laptop's <http://localhost> (<http://192.168.99.100> for me)
- In your docker attempts, you were using the host network (leveraging the `-p` port mapping command) to try to connect the containers
  - But maybe that didn't work because you actually do not know what IPs the containers are running at
  - It works for me, because I run them all on a linux VM that I've created using docker toolbox (hardware virtualization)
  - If you use Docker Desktop and do not create a new VM, you use your own OS to host your containers (container virtualization)
  - So you have to create a docker network to allow them to live in the same city (network) so they can talk to each other directly (not through the host computer)
- The `-p` port mappings are not really required anymore!
  - Only used to test the `/testHealth` and `/testComms` endpoints!
  - We do need it for the front end though, to access the UI!



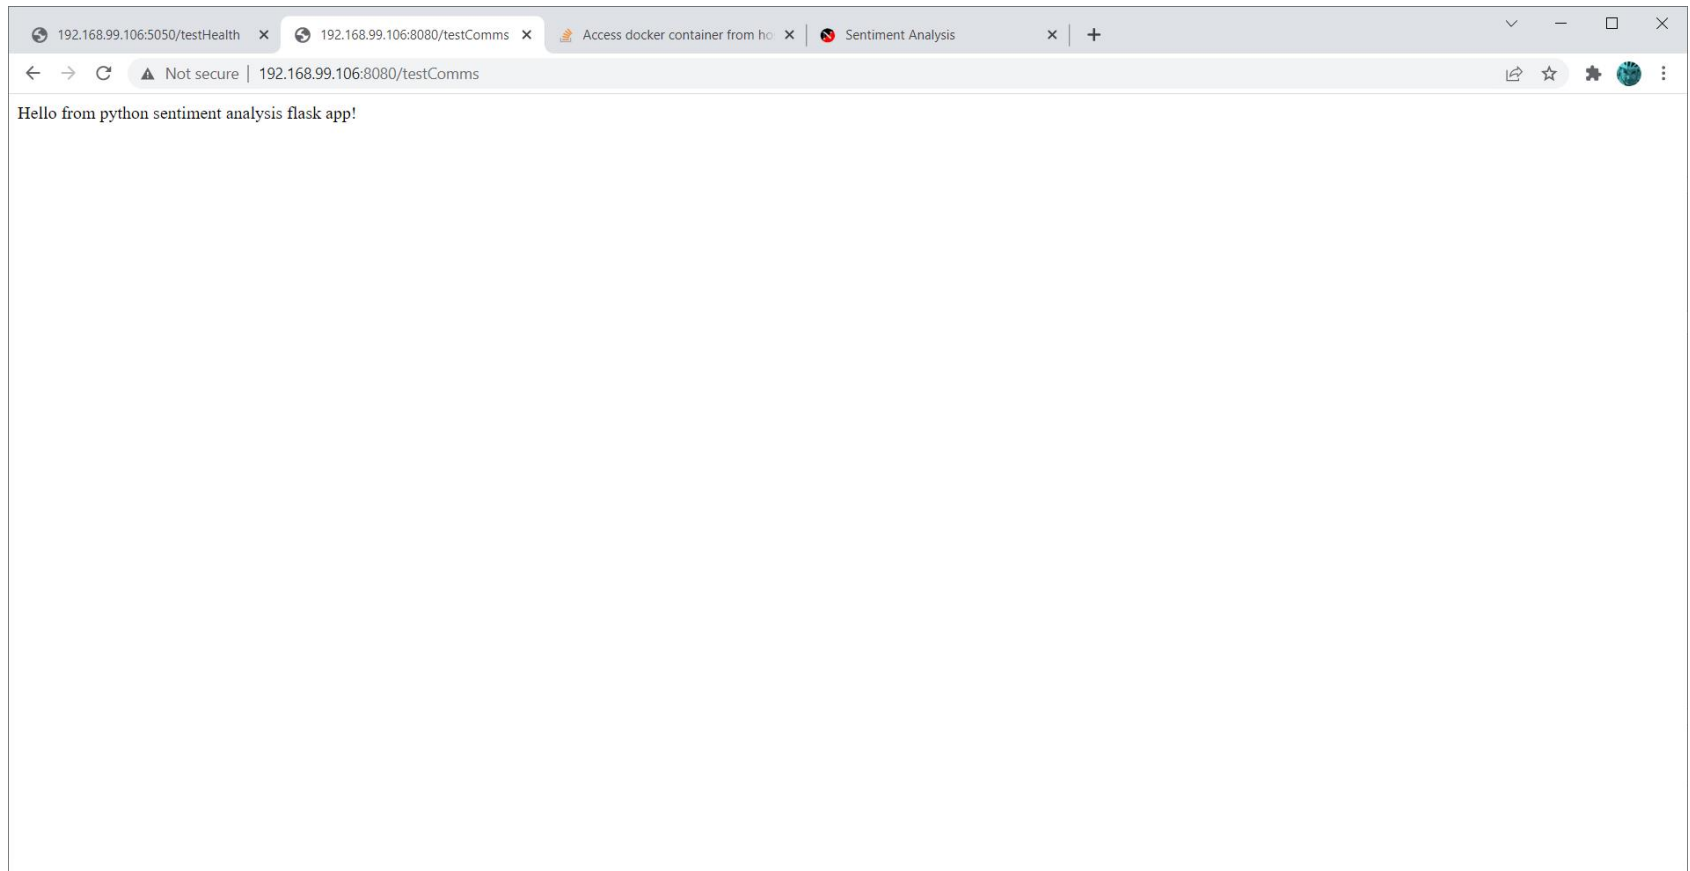


# /testHealth endpoints (192.168.99.106 host)

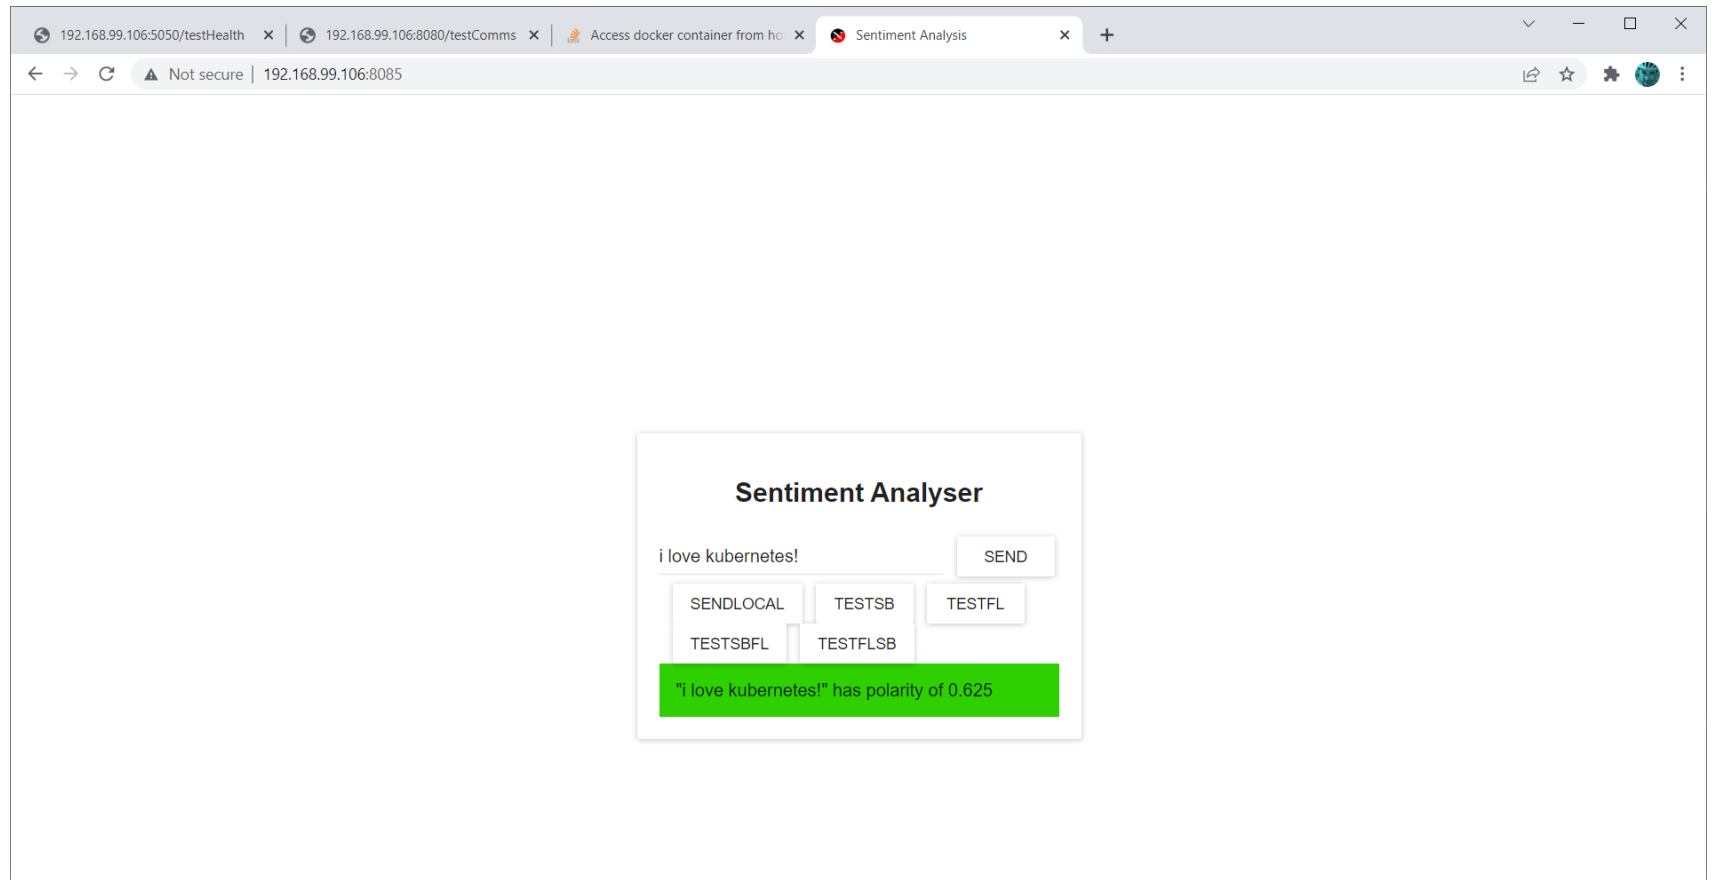




# /testComms endpoints (192.168.99.106 host)



# Front end





Part 2

# DOCKER-COMPOSE



# Docker Compose

- **Docker-compose** is a tool for defining and running multi-container Docker applications
- With Compose, we can create a YAML file to define all our microservices and with a single command, and spin everything up or tear it all down
- Make sure you have docker-compose in your installation
  - **docker-compose version**



# docker-compose

- With Compose, you use a YAML file to configure your application's services
- Then, with a single command, you create and start all the services from your configuration
- Also, **easy** to create **volumes** (map folders from host to container) and **networks** (to bridge containers)
- 3 steps:
  - **Write a Dockerfile for each  $\mu$ service**
  - **Define all  $\mu$ services that make up your ecosystem, in `docker-compose.yml`**
    - **Put that file outside of each service's folder**
  - **Run `docker compose up` and Docker starts and runs all your  $\mu$ services**
- <https://docs.docker.com/compose/>





# `docker-compose.yml` (outside of services folders)

```
#version: '3.8'
```

```
# see here https://docs.docker.com/compose/compose-file/compose-versioning/  
version: '3.1'
```

```
services:
```

```
  fe:
```

```
    build:
```

```
      context: ./sa-frontend
```

```
      dockerfile: Dockerfile
```

```
    ports:
```

```
      - 8085:80
```

```
    environment:
```

```
      - REACT_ENV=development
```

```
    networks:
```

```
      - frontend
```

**Name of service**

**Location of Dockerfile**

**Port mapping**

**Env vars**

**Creates a network**



# `docker-compose.yml` (continued)

#continuing

```
lo:
  build:
    context: ./sa-logic
    dockerfile: Dockerfile
  ports:
    - 8080:8080
  environment:
    - FLASK_ENV=development
  networks:
    - backend
```

**Name of service**

**Name of service**

```
wa:
  build:
    context: ./sa-webapp
    dockerfile: Dockerfile
  ports:
    - 8080:8080
  environment:
    - SA_LOGIC_API_URL=http://localhost:5000
    - WA_NETWORK=backend
  depends_on:
```

**Network specified  
as env var so that  
source code could  
use it**

```
  - lo
# specifies frontend and backend as the networks the wa service will have access to
networks:
  - frontend
  - backend
```

**networks  
that service wa  
has access to**



# `docker-compose.yml` (finish)

#continuing

# bridge networks to allow the containers to communicate with each other

networks:

frontend:

driver: bridge

backend:

driver: bridge

**Creates  
networks  
to bridge  
containers**



# Comment out env vars from the Dockerfile

```
FROM openjdk:8-jdk-alpine
# Environment Variable that defines the endpoint of sentiment-analysis python api:
#ENV SA_LOGIC_API_URL http://localhost:5000
ADD target/sentiment-analysis-web-0.0.1-SNAPSHOT.jar /
EXPOSE 8080
CMD ["java", "-jar", "sentiment-analysis-web-0.0.1-SNAPSHOT.jar", "--sa.logic.api.url=${SA_LOGIC_API_URL}"]
```



# Build $\mu$ service binaries first!

- **sa-logic:**
  - Nothing to do!
- **sa-webapp:**
  - `mvn install`
- **sa-frontend:**
  - `npm install`
  - `npm run build`

# docker-compose in action: Build step

- Build *all container* images in one step!
  - `docker-compose build`







# docker-compose in action: Run step

- Run the containers once the build is complete:
  - `docker-compose up -d`
- `docker container ls`
- Navigate to <http://localhost:8085> in your browser
  - (or <http://192.168.99.100:8085>)



# Docker teardown

- ❑ **To stop the containers:**
  - `docker-compose stop`
- ❑ **To bring down the containers:**
  - `docker-compose down`
- ❑ **Want to force a build?**
  - `docker-compose build --no-cache`
- ❑ **Remove images:**
  - `docker rmi $(docker images -q)`

An illustration of a magnifying glass with an orange handle. The lens is focused on a blue rectangular area containing several horizontal lines of varying lengths in white, green, and orange, representing code. A small orange bug with black legs and antennae is positioned on one of the lines. A thin, wavy line extends from the right side of the magnifying glass.

## Part 3 DEBUGGING CONTAINERS



# Debugging containers

- The docker exec command will let you run arbitrary commands inside an existing container
  - `docker exec -it <container_name> bash`
  - `docker exec -it <container_name> sh`
  - `docker exec -it <container_name> /bin/sh`
- To get into the container interactively
  - `docker run -it --entrypoint /bin/bash <container_name>`



# More on debugging containers

## □ View stdout history with the logs command

- `docker logs <container_name>`
- This history is available even after the container exits, as long as its file system is still present on disk (until it is removed with **docker rm**). The data is stored in a json file buried under `/var/lib/docker`

## □ Stream stdout with the attach command

- `docker attach <container_name>`
- By default this command attaches stdin and proxies signals to the remote process. Options are available to control both of these behaviors. To detach from the process use the default `ctrl-p ctrl-q` sequence

## □ Execute arbitrary commands with exec

- `docker exec <container_name> cat /var/log/test.log`

## □ interactive shell in the container

- `docker exec -it <container_name> /bin/sh`



# Extreme: Debug a container from another

- Create a debug container with strace
  - FROM alpine
  - RUN apk update && apk add strace
  - CMD ["strace", "-p", "1"]
- Build the container
  - docker build -t strace .
- Run strace container in the same pid and network namespace
  - docker run -t --pid=container:baadf00d \
  - --net=container: baadf00d \
  - --cap-add sys\_admin \
  - --cap-add sys\_ptrace \
  - strace
  - This attached strace to the baadf00d process and follows it as it executes





# Root filesystem

- To get to the root filesystem of the remote container, use the alpine image and launch a shell, in the same pid and network namespace
  - `docker run -it --pid=container:baadf00d \`  
  `--net=container:baadf00d \`  
  `--cap-add sys_admin \`  
  `alpine sh`
- **With this container attached to the original we can do more debugging**
  - You can still debug the network but make sure you use localhost because your new sh process is running in the same network namespace
  - `apk update && apk add curl lsof`  
  `curl localhost:2015`
  - `lsof -i TCP`
  - All standard debugging tools should work from this 2nd container without tainting the original container

